

SUMMARY

Data Engineer with 4+ years of experience designing and deploying scalable data pipelines and analytics solutions across AWS, Azure, and Snowflake environments. Proficient in SQL, Python, PySpark, and tools such as dbt, Airflow, S3, Glue, Azure Data Factory, Apache Kafka, and Terraform. Skilled in building dimensional data models (Star/Snowflake schemas), implementing data quality frameworks, optimizing Snowflake query performance and cloud costs, and developing Tableau and Power BI dashboards to drive business insights. Proven ability to troubleshoot ETL/ELT workflows, implement CI/CD pipelines, and support cross-functional teams in delivering timely, reliable data for reporting and analytics.

TECHNICAL SKILLS

Programming & Query Languages: Python, PySpark, SQL, Java, Bash

Cloud Platforms: AWS (S3, Redshift, Glue, Lambda), Azure (Data Factory, Synapse Analytics, Data Lake Storage, Functions)

Data Engineering, ETL, Orchestration & Reporting: Apache Airflow, Apache Kafka, Databricks, Delta Lake, dbt, Snowflake, Power BI, Tableau

Data Modeling & Transformation: Dimensional Modeling (Star & Snowflake Schemas), Schema Design & Evolution

Databases & Data Stores: PostgreSQL, MySQL, SQL Server, Oracle, DynamoDB

DevOps, Containers & Infrastructure: Git, Azure DevOps, Docker, Kubernetes, Terraform

Data Architecture, Quality, Governance & Compliance: Data Lake Architecture, Data Quality & Validation, Data Lineage, Pipeline Monitoring & Observability, Regulatory Compliance (GDPR, CCPA), Data Security & Access Control

AI/ML & GenAI: Scikit-learn, Pandas, NumPy, LangChain, RAG, Feature Engineering, Prompt Engineering

Methodologies & Soft Skills: Agile/Scrum, Leadership, Problem-Solving, Collaboration, Communication, Decision-Making

PROFESSIONAL EXPERIENCE

Data Engineer | Inspira Financial, USA (Remote) **Aug 2024 – Present**

- Developed PySpark ETL pipelines using AWS Glue and S3, ingesting 50M+ monthly transactions into Warehouses for analytics and DynamoDB for operational lookups, improving reporting efficiency, reducing processing time.
- Automated dbt transformations on Snowflake, producing 200+ governed datasets with model testing and run visibility, reducing manual intervention and accelerating analytics delivery cycles by 25%, enabling timely financial and risk reporting across business teams.
- Implemented reconciliation and audit workflows across portfolio and credit-risk data sources, standardizing controls and validations to ensure data accuracy, traceability, and compliant financial reporting in support of GDPR/CCPA requirements.
- Contributed to GenAI initiatives enabling governed exploration of retirement and financial data through conversational queries and automated insights, reducing ad-hoc analytics requests by 25% and accelerating self-service reporting across operations and risk teams.
- Designed Power BI dashboards leveraging curated Snowflake datasets, presenting portfolio KPIs, compliance metrics, and operational summaries, reducing reporting cycle time by 40% and improving executive visibility for stakeholders.
- Deployed containerized ML inference services on Kubernetes within AWS, using Terraform for infrastructure provisioning, improving model scalability and reliability for fraud analytics.

Data and Reporting Analyst | University of Massachusetts Dartmouth, Massachusetts, USA **Aug 2023 – July 2024**

- Extracted and transformed academic datasets using PostgreSQL and SQL, delivering 50+ institutional performance reports per semester, improving administrative decision-making accuracy and reducing data preparation time.
- Conducted thorough data validation and quality checks on datasets weekly using SQL, improving institutional reporting accuracy by 25%, ensuring reliable metrics for academic and administrative stakeholders.
- Applied Python for preprocessing and analysis, performing classification and clustering on student data to identify performance cohorts, enabling targeted interventions, curriculum adjustments, and supporting strategic faculty and administrative planning.
- Developed Tableau dashboards to visualize student enrollment trends, faculty performance, and course analytics, streamlining reporting workflows, reducing report preparation time, and enhancing leadership insights.
- Presented actionable insights to administration teams, enabling data-driven planning, optimizing course scheduling and resource allocation, and improving operational efficiency by 15% while supporting strategic academic initiatives.

Data Engineer | MTX Group, Hyderabad, India **Jan 2021 – Dec 2022**

- Orchestrated Apache Airflow pipelines integrated with Azure Data Factory, ingesting 100+ enterprise client datasets weekly, reducing processing delays by 20%, and meeting stringent SLAs for analytics and reporting deliverables.
- Transformed and optimized large-scale data using Databricks and Delta Lake, processing over 5TB per month, implementing incremental ingestion and schema evolution for client reporting, analytics, and operational insights.
- Curated and structured datasets in SQL Server and Oracle, producing operational reports for cross-functional teams, enhancing audit readiness, improving reporting accuracy, and increasing decision-making efficiency.
- Managed Git version control and automated Azure DevOps CI/CD pipelines, reducing manual errors, accelerating sprint delivery by 18%, and ensuring consistent deployment of data workflows under Agile/Scrum methodology.
- Collaborated with architects, QA, and data teams to validate outputs in Azure Data Lake Storage and Azure Synapse Analytics, improving data accuracy, regulatory compliance, and reliability of client dashboards.
- Investigated and resolved pipeline failures using Azure Monitor and Airflow logs, reducing downtime by 40%, improving data availability, and ensuring uninterrupted delivery of client reports and analytical insights.

Data Engineer Intern Zemoso Technologies , Hyderabad, India	July 2020 – Dec 2020
- Assisted in AWS Lambda ingestion workflows with S3, processing 500K+ records per month, contributing to ETL delivery and helping ensure reliability of downstream analytics pipelines. - Executed feature engineering using NumPy on dataset variables, preparing clean inputs for machine learning models, internal reporting, and downstream analytics pipelines, improving dataset usability. - Performed data quality checks on ETL jobs weekly, reducing data errors by 20%, enhancing accuracy of analytics outputs, and providing reliable datasets for senior engineers and production teams.	

CERTIFICATIONS

- SnowPro Core Certification
- dbt Fundamentals Badge – dbt Labs
- Databricks Fundamentals Accreditation
- Snowflake Hands-on Labs (Project Badges) – Snowflake
- Generative AI with Large Language Models – DeepLearning.AI (Coursera)

EDUCATION

University of Massachusetts Dartmouth , Massachusetts, USA	Jan 2023 – Dec 2024
Master of Science in Data Science	
Vignana Bharathi Institute of Technology , Hyderabad, INDIA	Aug 2017 – Aug 2021
Bachelor of Technology in Information Technology	